

SMART BUDGETING AGENT

A full-stack AI system that automates financial transaction categorization

Derrick K. Gibbs-McGlaston

Developer and designer of the complete system

Capstone ITAI-2277

Final project for the AI Programming Capstone course

Instructor: Sitaram Ayyagari

Problem Statement

Why budgeting is hard

People struggle because expense tracking
is tedious and inconsistent

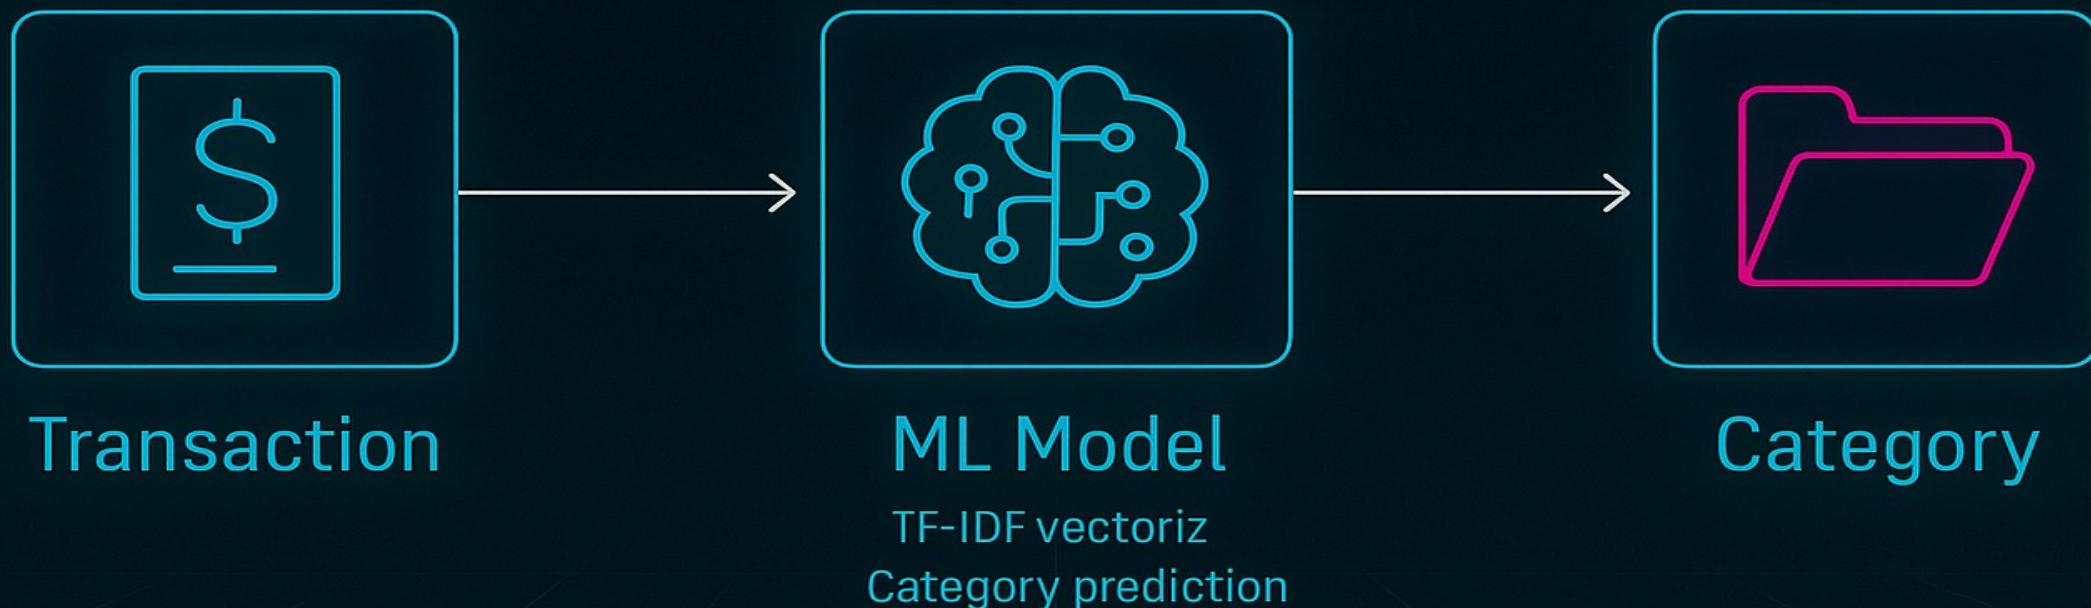
Why categorization matters

Clean categories help users understand
spending patterns and improve financial
decisions

What gap your system solves

It removes manual work by predicting
categories instantly using machine learning

EXPENSE CATEGORIZATION



Project Goals

Automate expense categorization

Replace manual labeling with a fast, accurate ML model

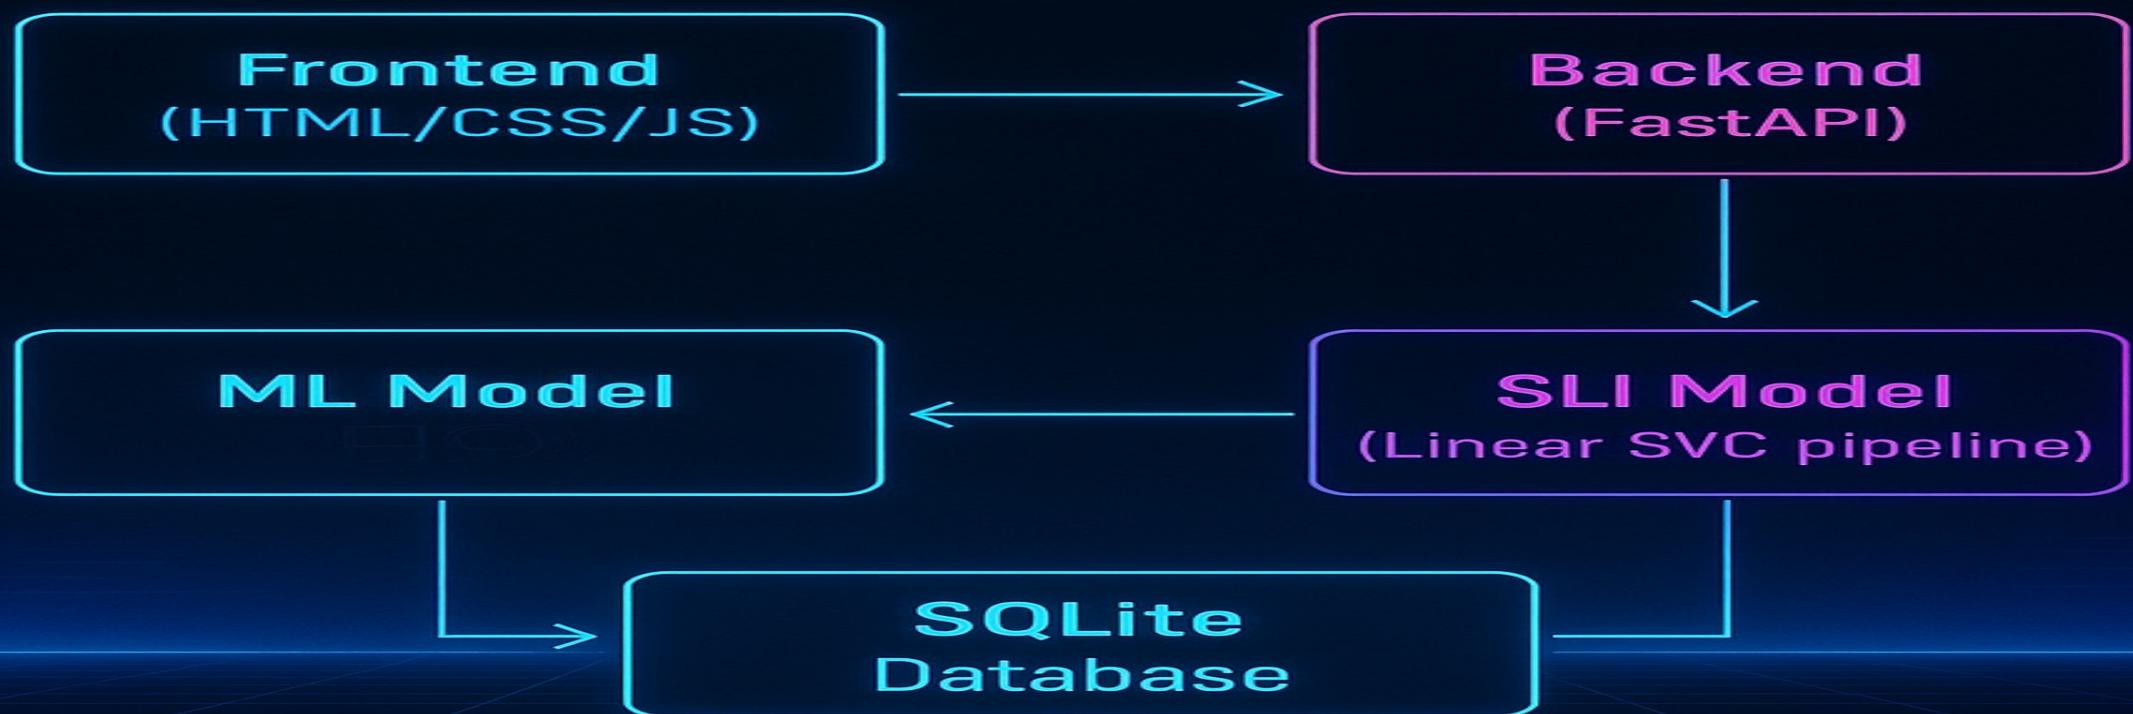
Use ML + full-stack design

Integrate machine learning with a real API and a functional UI

Build a functional deployable system

Deliver a prototype that can be installed, run, and tested end-to-end

System Architecture



DATA PIPELINE

SOURCE OF DATA

Transactions collected and labeled for training

CLEANING

Removed noise, punctuation, special characters, and normalized text

TOKENIZATION

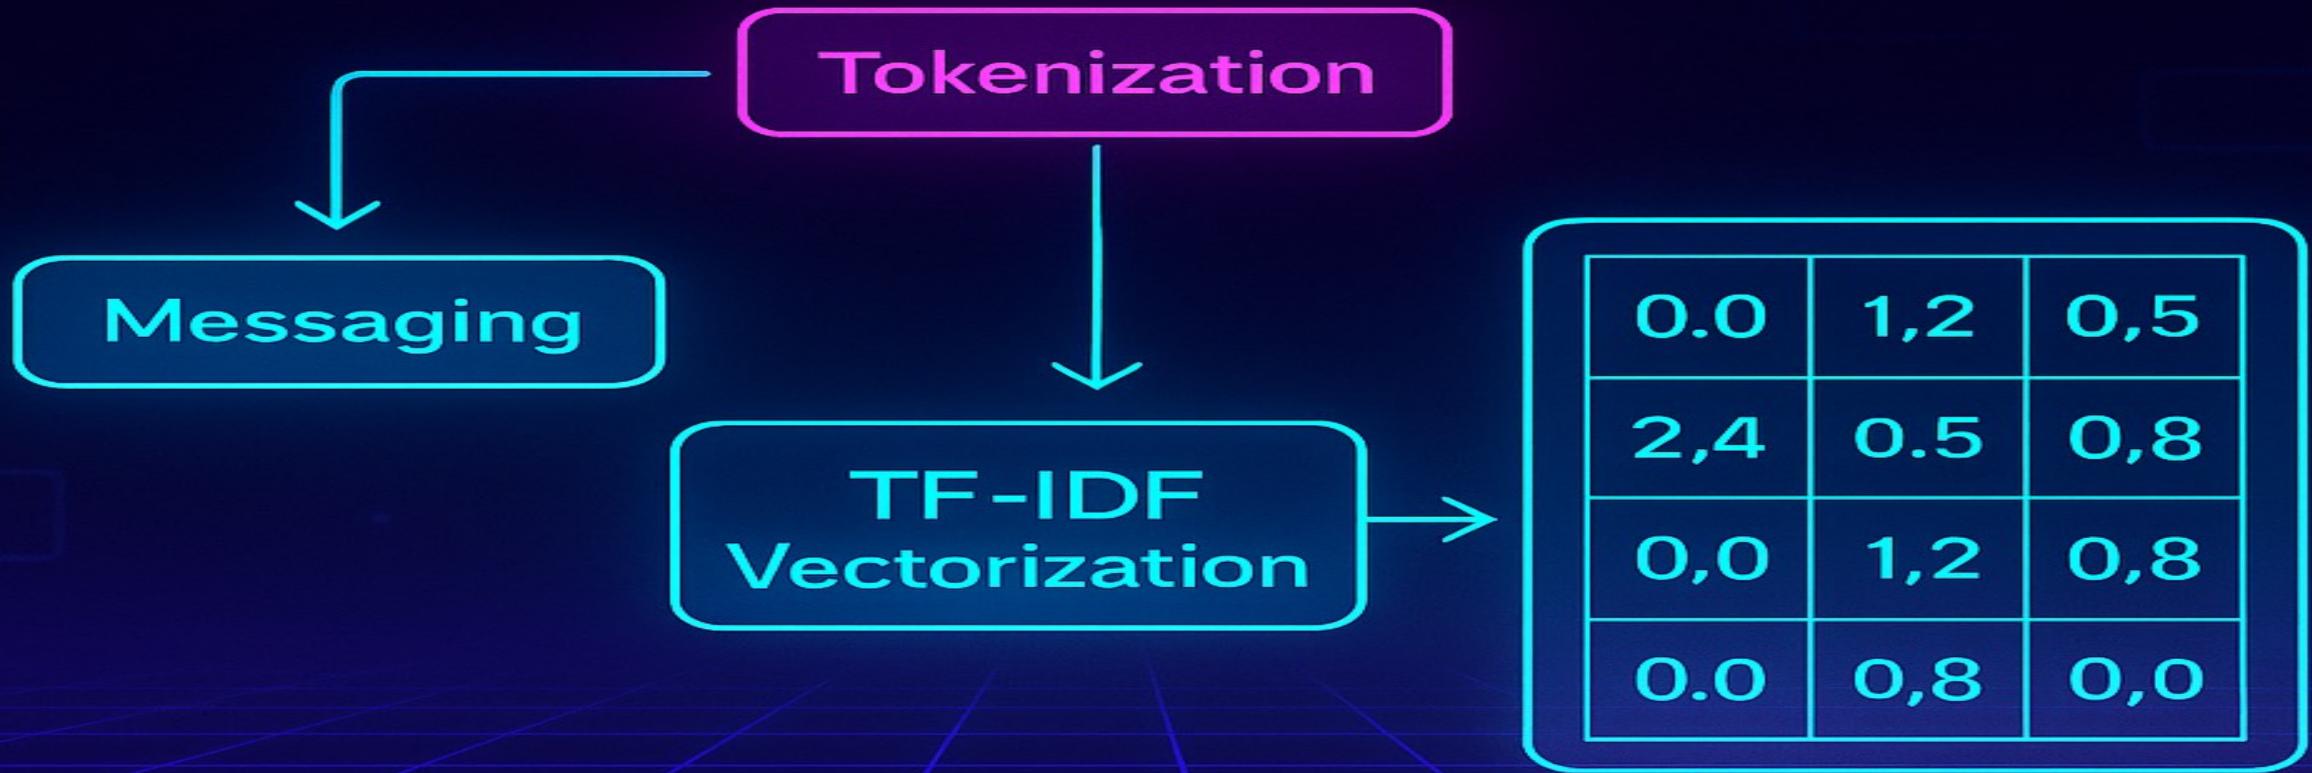
Broke text into meaningful units for ML processing

VECTORIZATION

Converted words into TF-IDF numerical features

TRAIN/TEST SPLIT

Ensured unbiased evaluation of model performance



MACHINE LEARNING MODEL

- **LinearSVC** – Chosen for speed and strong performance on text classification tasks
- **TF-IDF** – Transforms text into weighted numerical vectors the model can learn from
- **Accuracy** – Achieved approx. 90% accuracy on test data
- **Why SVC fits the problem** – It handles sparse text data extremely well with low computational cost



BACKEND API DESIGN

- `/predict` – Accepts transaction data and returns a predicted category
- `TF-IDF` – Transforms text into weighted numerical structure for clean storage
- `Schema.sql` – Defines the SQLite database structure for clean storage
- Security + performance considerations – Lightweight server, isolated

REST Pipeline



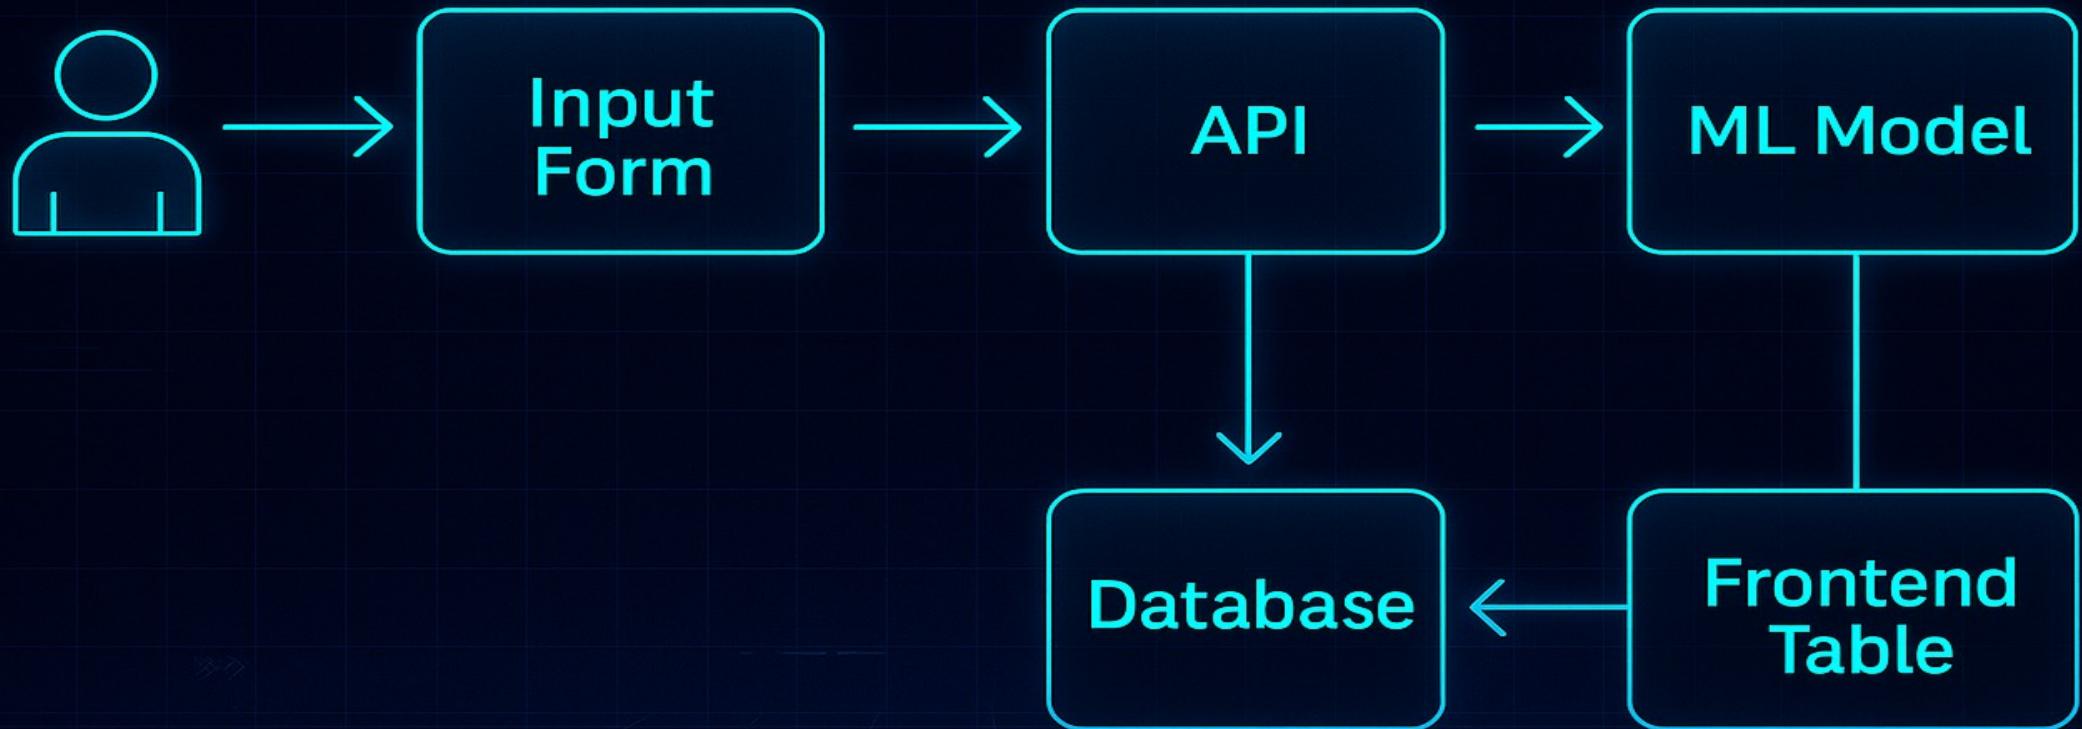
FRONTEND INTEGRATION

- Prediction form – Users enter transaction details quickly and easily
- Live updates – Predictions show instantly without reloading the page
- Transaction table – Displays stored transactions pulled from backend
- UX simplicity – Clean, minimal design focused on functionality over visuals

DEMO FLOW

- Enter transaction – User fills in description, merchant, amount, etc.
- Save prediction – Backend returns the category predicted by the ML model
- Save to DB – The transaction and prediction are stored in SQLite
- View in recent table – Updated record appears instantly in the frontend

PREDICTION FLOW





TECHNICAL CHALLENGES

- Environment conflicts – Issues with virtual environments and dependency mismatches
- Broken model predictions – Initial models produced inconsistent results
- Fixing label quality – Improved training data for better accuracy
- FastAPI debugging – Resolved server errors and routing issues

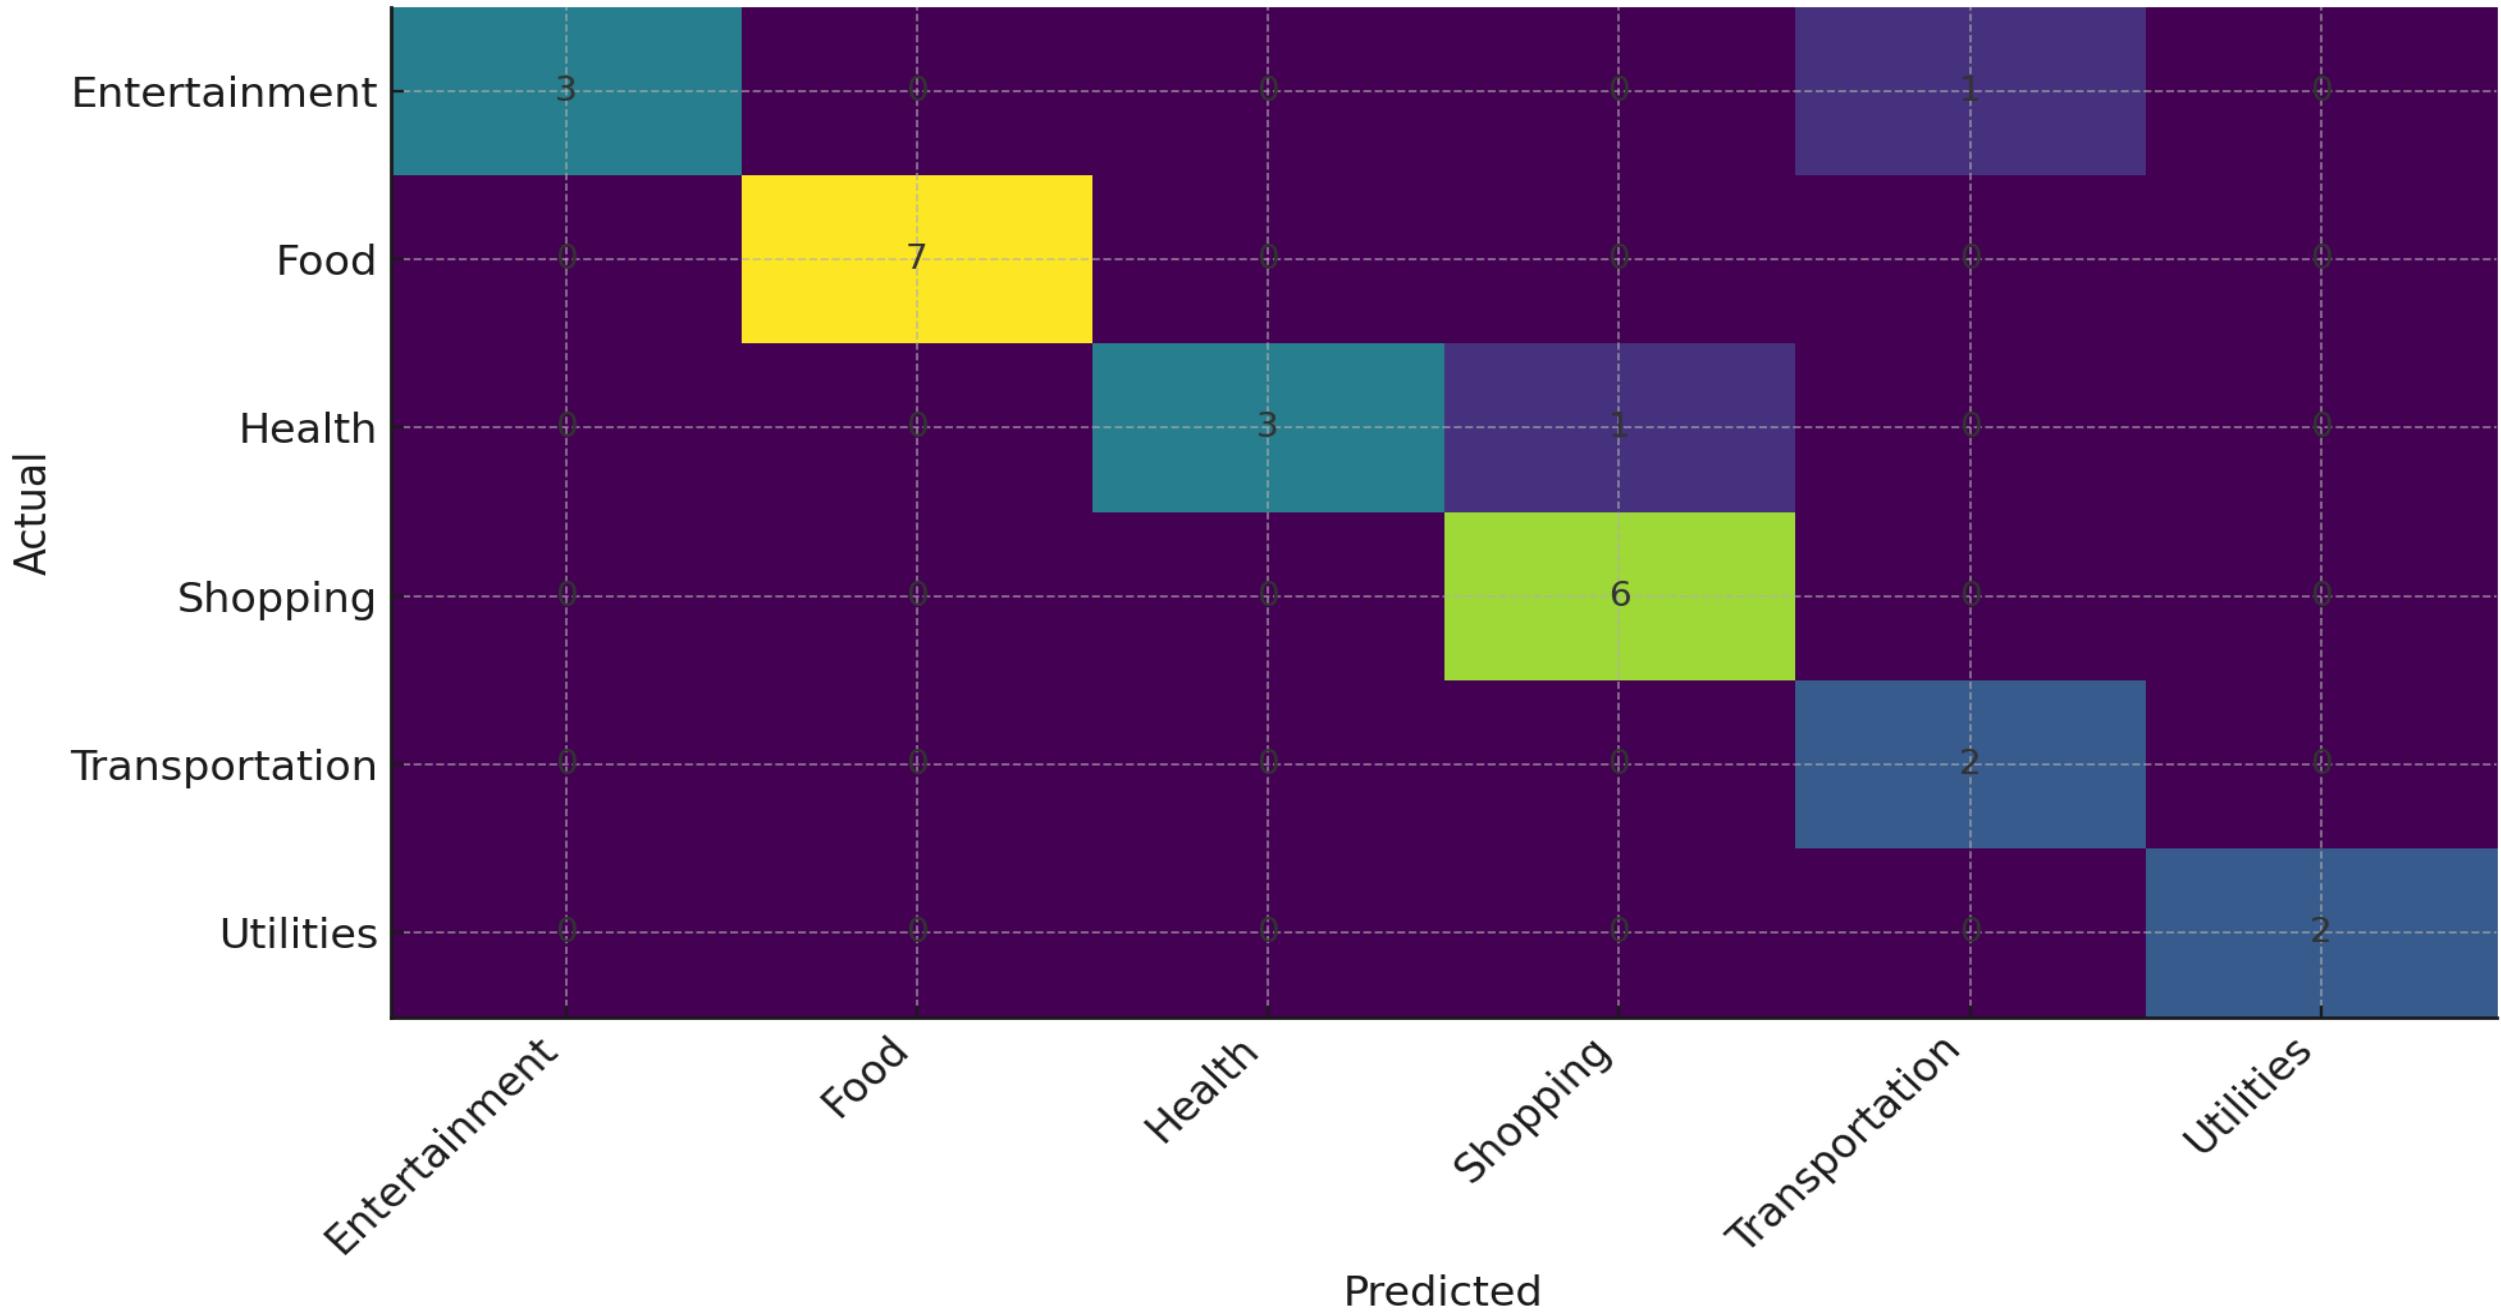
SOLUTIONS IMPLEMENTED

- **Rebuilt model pipeline** – Created a stable repeatable training workflow
- **Added override logic** – Handled cases where certain merchants map to fixed categories
- **Cleaned virtual env** – Removed conflicts and reinstalled dependencies cleanly
- **Improved training** – Refined labels and feature engineering

EVALUATION METRICS

- **Confusion matrix** – Visual breakdown of model performance by class
- **Added override logic** – Precision, recall, and F1 scores for each category
- **Accuracy** – Demonstrated strong overall predictive performance
- **Error patterns** – Identified which categories were most difficult to classify

Decision Tree Confusion Matrix (Test)



CAPSTONE REQUIREMENTS

- ML model – Complete ML pipeline with documentation
- Full stack integration – Precision, recall, and F1 scores for each category
- Repo + documentation – Clean GitHub repository with full instructions
- Demo-ready prototype – Fully operational and testable

REAL-WORLD USE CASES

- Personal finance – Individuals wanting clearer budgeting insights
- Fin-tech budgeting apps – Foundations for consumer-level financial tools
- Banking transactions – Automated categorization systems used in mobile banking
- Foundations for FussBudget – Can evolve into a full budgeting platform

FUTURE ENHANCEMENTS

- User accounts – Add authentication and personalized dashboards
- Cloud deployment – Host backend and frontend online for public use
- Multiple models – Compare SVC with neural models for improved accuracy
- Budget alerts – Notify users when they overspend

ETHICAL CONSIDERATIONS

- Data privacy – Ensure user data is protected and never misused
- Model bias – Monitor for biased predictions based on merchant or text patterns
- Transparency – Explain model decisions where possible
- Security considerations – Protect API endpoints and database integrity

LESSONS LEARNED

- ML pipeline structuring – Understanding how to properly build evaluate and save models
- Debugging real systems – Handling errors across multiple layers of a stack
- End-to-end development – Connecting data, model, backend, and frontend
- GitHub professionalisms – Publishing clean, testable, employer-ready repos

CONCLUSION

- Restate what you built – A complete AI-powered expense categorization system
- Impact of the system – Makes budgeting easier, faster, and more accurate
- Skills demonstrated – Machine learning full-stack development, and technical problem-solving

THANK YOU